

# An example of medical treatment optimization under model uncertainty

Orlane Le Quellenec <sup>1</sup>, Alice Cleyne <sup>1,2</sup>, Benoîte de Saporta <sup>1</sup> and Régis Sabbadin <sup>3</sup>

<sup>1</sup>IMAG, Univ Montpellier, CNRS, Montpellier, France

<sup>2</sup>John Curtin School of Medical Research, The Australian National University, Canberra, ACT, Australia

<sup>3</sup>Univ Toulouse, INRAE-MIAT, Toulouse, France

June 30, 2023



UNIVERSITÉ DE  
MONTPELLIER

INRAE

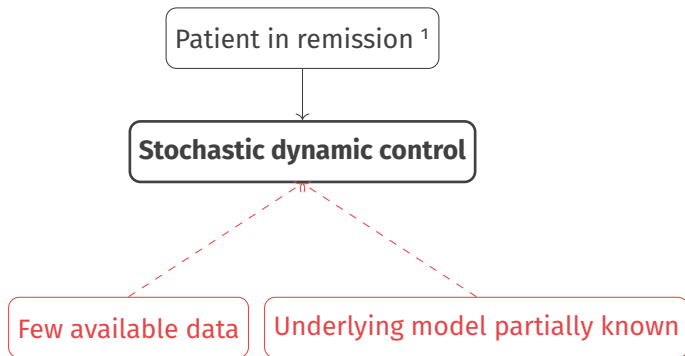
IMAG

INSTITUT MONTPELLIERAIN  
ALEXANDER GROTHENDIECK



anr<sup>®</sup>

# A medical context

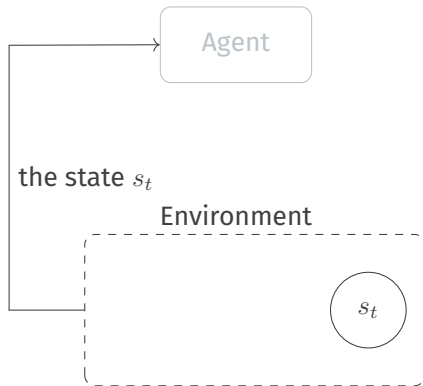


**How can these issues be addressed in a simplified problem ?**

---

<sup>1</sup>Data from IUC Oncopole, Toulouse, and CRCT, Toulouse, France

# Markov Decision Process (MDP<sup>2</sup>)

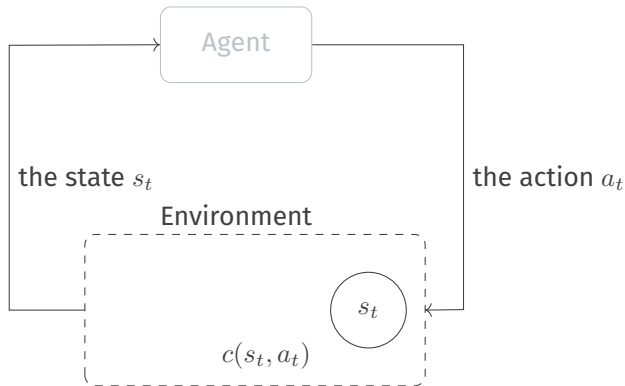


- $s \in \mathcal{S}$  the state space
- $a \in \mathcal{A}$  the action space
- $\mathcal{P}$  the transition matrix
- $c(s_t, a_t)$  the cost function

---

<sup>2</sup>ML Puterman (1994). "Finite-horizon Markov decision processes". In: Markov Decision Processes: Discrete Stochastic Dynamic Programming. New York: Wiley-Interscience, pp. 78–9.

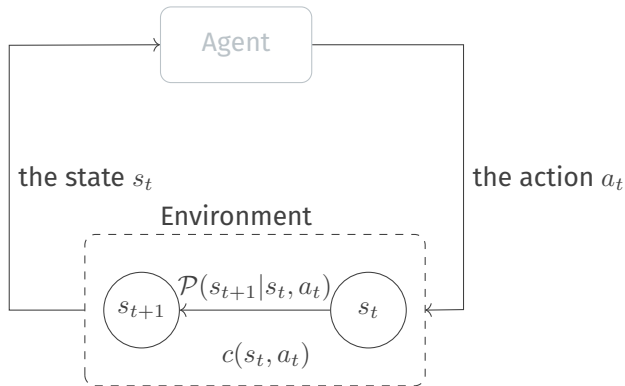
# Markov Decision Process (MDP<sup>2</sup>)



- $s \in \mathcal{S}$  the state space
- $a \in \mathcal{A}$  the action space
- $\mathcal{P}$  the transition matrix
- $c(s_t, a_t)$  the cost function

<sup>2</sup>ML Puterman (1994). "Finite-horizon Markov decision processes". In: Markov Decision Processes: Discrete Stochastic Dynamic Programming. New York: Wiley-Interscience, pp. 78–9.

# Markov Decision Process (MDP<sup>2</sup>)



- $s \in \mathcal{S}$  the state space
- $a \in \mathcal{A}$  the action space
- $\mathcal{P}$  the transition matrix
- $c(s_t, a_t)$  the cost function

<sup>2</sup>ML Puterman (1994). "Finite-horizon Markov decision processes". In: Markov Decision Processes: Discrete Stochastic Dynamic Programming. New York: Wiley-Interscience, pp. 78–9.

# State transition

The transition matrix is partially known

**Table: Transition matrix when patient has no treatment ( $a = \emptyset$ ).**

$s_t \backslash s_{t+1}$	(0,0,0)	(1,0,1)	(1,0,2)	(1,1,1)	(1,1,2)	(1,2,1)	(1,2,2)	(1,3,1)	(1,3,2)	(2,4,0)
(0,0,0)	$p_{(0,0,0)}^\emptyset$	$p_{(1,0,1)}^\emptyset$	$p_{(1,0,2)}^\emptyset$	0	0	0	0	0	0	0
(1,0,1)	0	0	0	<b>1</b>	0	0	0	0	0	0
(1,0,2)	0	0	0	0	0	0	<b>1</b>	0	0	0
(1,1,1)	0	0	0	0	0	<b>1</b>	0	0	0	0
(1,1,2)	0	0	0	0	0	0	0	0	<b>1</b>	0
(1,2,1)	0	0	0	0	0	0	0	<b>1</b>	0	0
(1,2,2)	0	0	0	0	0	0	0	0	0	<b>1</b>
(1,3,1)	0	0	0	0	0	0	0	0	0	<b>1</b>
(1,3,2)	0	0	0	0	0	0	0	0	0	<b>1</b>
(2,4,0)	0	0	0	0	0	0	0	0	0	<b>1</b>

# State transition

The transition matrix is partially known

Table: Transition matrix when patient has treatment ( $a = \rho$ ).

$s_t \backslash s_{t+1}$	(0,0,0)	(1,0,1)	(1,0,2)	(1,1,1)	(1,1,2)	(1,2,1)	(1,2,2)	(1,3,1)	(1,3,2)	(2,4,0)
(0,0,0)	$p_{(0,0,0)}^\rho$	$p_{(1,0,1)}^\rho$	$p_{(1,0,2)}^\rho$	0	0	0	0	0	0	0
(1,0,1)	1	0	0	0	0	0	0	0	0	0
(1,0,2)	1	0	0	0	0	0	0	0	0	0
(1,1,1)	1	0	0	0	0	0	0	0	0	0
(1,1,2)	1	0	0	0	0	0	0	0	0	0
(1,2,1)	0	0	0	1	0	0	0	0	0	0
(1,2,2)	0	0	0	0	1	0	0	0	0	0
(1,3,1)	0	0	0	0	0	0	0	1	0	0
(1,3,2)	0	0	0	0	0	0	0	0	1	0
(2,4,0)	0	0	0	0	0	0	0	0	0	1

# Solving a MDP

Minimizing a cost

The list of costs:

- Treatment: 300
- Disease 1: 200
- Disease 2: 300
- Death: 1000

## Policy $\pi$

Let  $f : \mathcal{S} \rightarrow \mathcal{A}$  for all  $s \in \mathcal{S}$  is a decision rule.

A sequence of decision rules  $\pi = (f_0, f_1, \dots, f_{H-1})$  is a policy.

## Policy cost

$$J_H(\pi, s) = \mathbb{E}\left[\sum_{t=0}^{H-1} c(s_t, a_t) | \pi, s\right]$$



# Solving a MDP

Minimizing a cost

The list of costs:

- Treatment: 300
- Disease 1: 200
- Disease 2: 300
- Death: 1000

## Policy $\pi$

Let  $f : \mathcal{S} \rightarrow \mathcal{A}$  for all  $s \in \mathcal{S}$  is a decision rule.

A sequence of decision rules  $\pi = (f_0, f_1, \dots, f_{H-1})$  is a policy.

## Policy cost

$$J_H(\pi, s) = \mathbb{E}\left[\sum_{t=0}^{H-1} c(s_t, a_t) \mid \pi, s\right]$$

## Optimization criterion

$$V^*(s_t) = \min_{a \in \mathcal{A}} [c(s_t, a) + \sum_{s_{t+1} \in \mathcal{S}} \mathcal{P}(s_{t+1} \mid s_t, a) V^*(s_{t+1})]$$

# A model-free method

Q-learning<sup>3,4</sup> algorithm




---

<sup>3</sup>Christopher J. C. H. Watkins and Peter Dayan (May 1992). “Q-learning”. In: Mach. Learn. 8.3, pp. 279–292. ISSN: 1573-0565. DOI: [10.1007/BF00992698](https://doi.org/10.1007/BF00992698).

<sup>4</sup>VP Vivek and Dr. Shalabh Bhatnagar (Aug. 2022). “Finite Horizon Q-learning: Stability, Convergence, Simulations and an application on Smart Grids”. In: arXiv:2110.15093v3. DOI: [10.48550/arXiv.2110.15093](https://doi.org/10.48550/arXiv.2110.15093). eprint: [2110.15093v3](https://arxiv.org/abs/2110.15093v3).

# A model-free method

Q-learning<sup>3,4</sup> algorithm



---


<sup>3</sup>Christopher J. C. H. Watkins and Peter Dayan (May 1992). “Q-learning”. In: Mach. Learn. 8.3, pp. 279–292. ISSN: 1573-0565. DOI: [10.1007/BF00992698](https://doi.org/10.1007/BF00992698).

<sup>4</sup>VP Vivek and Dr. Shalabh Bhatnagar (Aug. 2022). “Finite Horizon Q-learning: Stability, Convergence, Simulations and an application on Smart Grids”. In: arXiv:2110.15093v3. DOI: [10.48550/arXiv.2110.15093](https://doi.org/10.48550/arXiv.2110.15093). eprint: [2110.15093v3](https://arxiv.org/abs/2110.15093v3).

# A model-free method

Q-learning<sup>3,4</sup> algorithm

$$Q_n(s_t, a_t) = (1 - \alpha)Q_{n-1}(s_t, a_t) + \alpha[c(s_t, a_t) + \min_{a_{t+1} \in \mathcal{A}} Q_{n-1}(s_{t+1}, a_{t+1})]$$




<sup>3</sup>Christopher J. C. H. Watkins and Peter Dayan (May 1992). “Q-learning”. In: Mach. Learn. 8.3, pp. 279–292. ISSN: 1573-0565. DOI: [10.1007/BF00992698](https://doi.org/10.1007/BF00992698).

<sup>4</sup>VP Vivek and Dr. Shalabh Bhatnagar (Aug. 2022). “Finite Horizon Q-learning: Stability, Convergence, Simulations and an application on Smart Grids”. In: arXiv:2110.15093v3. DOI: [10.48550/arXiv.2110.15093](https://doi.org/10.48550/arXiv.2110.15093). eprint: [2110.15093v3](https://arxiv.org/abs/2110.15093v3).

# A model-free method

Q-learning<sup>3,4</sup> algorithm

$$Q_n(s_t, a_t) = (1 - \alpha)Q_{n-1}(s_t, a_t) + \alpha[c(s_t, a_t) + \min_{a_{t+1} \in \mathcal{A}} Q_{n-1}(s_{t+1}, a_{t+1})]$$



<sup>3</sup>Christopher J. C. H. Watkins and Peter Dayan (May 1992). “Q-learning”. In: *Mach. Learn.* 8.3, pp. 279–292. ISSN: 1573-0565. DOI: [10.1007/BF00992698](https://doi.org/10.1007/BF00992698).

<sup>4</sup>VP Vivek and Dr. Shalabh Bhatnagar (Aug. 2022). “Finite Horizon Q-learning: Stability, Convergence, Simulations and an application on Smart Grids”. In: [arXiv:2110.15093v3](https://arxiv.org/abs/2110.15093v3). DOI: [10.48550/arXiv.2110.15093](https://doi.org/10.48550/arXiv.2110.15093). eprint: [2110.15093v3](https://arxiv.org/abs/2110.15093v3).

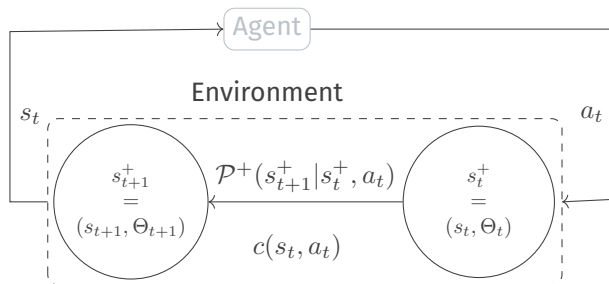
# A bayesian approach

$s_t \backslash s_{t+1}$	(0, 0, 0)	(1, 0, 1)	(1, 0, 2)	(1, 1, 1)	(1, 1, 2)	(1, 2, 1)	(1, 2, 2)	(1, 3, 1)	(1, 3, 2)	(2, 4, 0)
(0, 0, 0)	$p_{(0,0,0)}^\emptyset$	$p_{(1,0,1)}^\emptyset$	$p_{(1,0,2)}^\emptyset$	○	○	○	○	○	○	○

Remark:

- $P(.|s = (0, 0, 0), a = \emptyset) \sim \mathcal{M}(p_{(0,0,0)}^\emptyset, p_{(1,0,1)}^\emptyset, p_{(1,0,2)}^\emptyset)$
- Conjugate distribution :  $f(p^\emptyset | \Theta^\emptyset) \sim \mathcal{D}(\theta_{(0,0,0)}^\emptyset, \theta_{(1,0,1)}^\emptyset, \theta_{(1,0,2)}^\emptyset)$

# Bayes-Adaptive Markov Decision Process (BAMDP<sup>5</sup>)

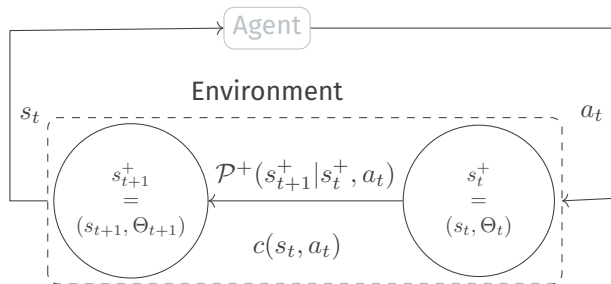


- $s^+ \in \mathcal{S}^+$  the hyper-state space
- $\mathcal{P}^+$  the transition matrix
- $\Theta_{t+1} = \Theta_t + \Delta_{s_{t+1}}^{a_t}$ , with

$$\Delta_{s_{t+1}}^{a_t} = \begin{cases} 1 & \text{if } (s = (\text{0}, \text{0}, \text{0}), a_t, s_{t+1}), \\ 0 & \text{else.} \end{cases}$$

<sup>5</sup>Michael O’Gordon Duff (2002). “Optimal learning: Computational procedures for Bayes -adaptive Markov decision processes”. PhD thesis. University of Massachusetts Amherst.

# Bayes-Adaptive Markov Decision Process (BAMDP<sup>5</sup>)



- $s^+ \in \mathcal{S}^+$  the hyper-state space
- $\mathcal{P}^+$  the transition matrix
- $\Theta_{t+1} = \Theta_t + \Delta_{s_{t+1}}^{a_t}$ , with

$$\Delta_{s_{t+1}}^{a_t} = \begin{cases} 1 & \text{if } (s = (\text{0}, \text{0}, \text{0}), a_t, s_{t+1}), \\ 0 & \text{else.} \end{cases}$$

Optimization criterion

$$V^*(s_t, \Theta_t) = \min_{a \in \mathcal{A}} [c(s_t, a_t) + \sum_{s_{t+1}^+ \in \mathcal{S}^+} \mathcal{P}^+(s_{t+1}^+ | s_t^+, a_t) V^*(s_{t+1}, \Theta_{t+1})]$$

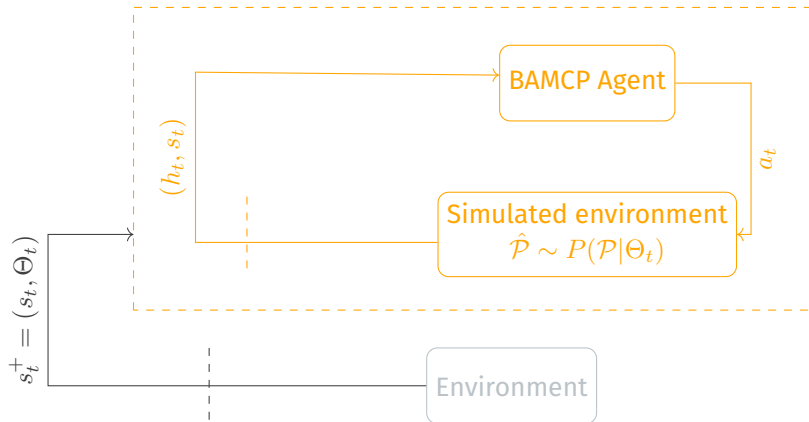
<sup>5</sup>Michael O’Gordon Duff (2002). “Optimal learning: Computational procedures for Bayes -adaptive Markov decision processes”. PhD thesis. University of Massachusetts Amherst.



# A model-based method

Bayes-Adaptive Monte-Carlo Planning (BAMCP)<sup>6</sup>

with  $h_t = (s_0, a_0, s_1, \dots, s_{t-1}, a_{t-1})$ ,

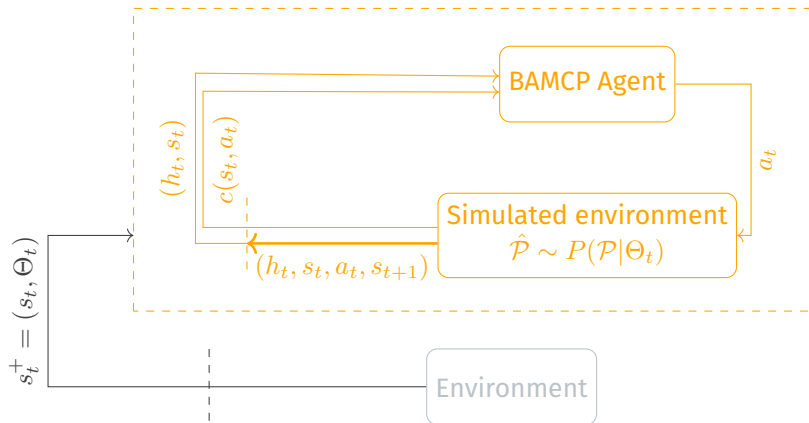


<sup>6</sup>Arthur Guez, David Silver, and Peter Dayan (2012). "Efficient Bayes-Adaptive Reinforcement Learning using Sample-Based Search". In: *Advances in Neural Information Processing Systems*. Ed. by F. Pereira et al. Vol. 25. Curran Associates, Inc.

# A model-based method

Bayes-Adaptive Monte-Carlo Planning (BAMCP)<sup>6</sup>

with  $h_t = (s_0, a_0, s_1, \dots, s_{t-1}, a_{t-1})$ ,

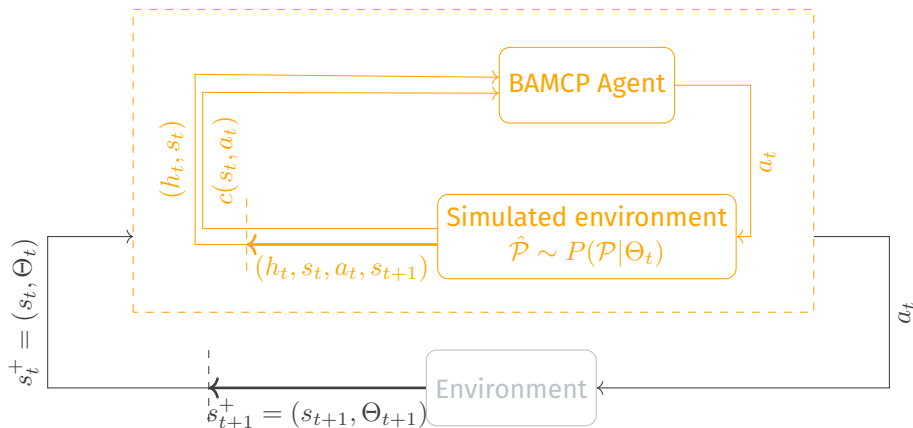


<sup>6</sup>Arthur Guez, David Silver, and Peter Dayan (2012). "Efficient Bayes-Adaptive Reinforcement Learning using Sample-Based Search". In: *Advances in Neural Information Processing Systems*. Ed. by F. Pereira et al. Vol. 25. Curran Associates, Inc.

# A model-based method

Bayes-Adaptive Monte-Carlo Planning (BAMCP)<sup>6</sup>

with  $h_t = (s_0, a_0, s_1, \dots, s_{t-1}, a_{t-1})$ ,




<sup>6</sup>Arthur Guez, David Silver, and Peter Dayan (2012). "Efficient Bayes-Adaptive Reinforcement Learning using Sample-Based Search". In: *Advances in Neural Information Processing Systems*. Ed. by F. Pereira et al. Vol. 25. Curran Associates, Inc.

# Results

The optimal policy exact cost: 888.89


Simulated patients	Q-learning		BAMCP	
	Cost	Time	Cost	Time
$10^2$	$1427.06 \pm 1.05$	0.07 sec	$1302.58 \pm 1.32$	2.07 hours
$10^3$	$936.96 \pm 0.70$	2.48 min	$1297.64 \pm 1.32$	2.22 hours
$10^4$	$936.93 \pm 0.70$	4.17 min	NC	4 days
$10^6$	$891.6 \pm 0.68$	10.21 min	NC	1.5 years

# Conclusion



● Mathematical framework    ● Model-free method    ● Model-based method

# Conclusion



- Mathematical framework
- ◆ Model-free method
- Model-based method

# Perspectives

MDP model	→	PDMP <sup>7</sup> model
Finite state space	→	Continuous state space
Markovian	→	Semi-Markovian
Complete observations	→	Hidden observations

Unlike model-free methods and deep reinforcement learning, **bayesian approaches** do not require as much interaction with the environment.

---

<sup>7</sup>Mark H. A. Davis (1984). “Piecewise-Deterministic Markov Processes: A General Class of Non-Diffusion Stochastic Models”. In: Journal of the Royal Statistical Society Series B (Methodological) 46, pp. 353–376.